

Group 6 Assignment Report

An Evolutionary Algorithm for solving Sudokus

Ben Bekir Ertugrul Eduard Iorga Haoran Liu
Moritz Schlitter

October 8, 2025

1 Introduction

The recommended software for typesetting assignment reports is L^AT_EX. It will allow you to prepare high-quality documents, especially in the area of Computer Science. This document can serve as a template for reports. Each section begins with brief instructions in red text. All the instructions in red, as well as the dummy text, should be removed in the final version to submit. The L^AT_EX source of this file includes examples of using the most needed commands and environments. You can find plenty of other examples with explanations in many web forums and discussion groups on the Internet. The easiest way to edit your report is to use <https://www.overleaf.com/>. Overleaf does not require any setup on your computer, and it is free to create an account.

The book *Writing for Computer Science* [?] is a useful assistance on how to write properly and present your work when it comes to Computer Science topics. It is a strong recommendation to follow its guidelines and limit the usage of AI tools to generate text. Keep in mind that the examiner is an expert in Evolutionary Computation and therefore, any false information generated by an AI tool is easily notable. Such case may lead to failing the assignment.

The introduction should briefly introduce the assignment and its purpose.

Traditionally, a Sudoku is a logic based number placement puzzle. The objective of the puzzle is to fill a most commonly 9x9 sized grid with digits so that each column, row and 3x3 subgrid contain all of the digits from 1 to 9.[?]

Solving such a puzzle programmatically falls into the category of search and optimization problems. These types of problems can be approached in different ways. One possible approach is depicted by evolutionary algorithms or in this case more specifically, genetic algorithms (GA)[1].

Inspired by the process of natural selection, GAs use biologically inspired operators such as selection, crossover and mutation to generate solutions to optimization and search problems[2].

This paper studies the efficiency of solving Sudoku puzzles with GA approaches. The objective is to compare the performance of different implementations of GAs. These approaches will additionally be compared to the naive search algorithm depth-first search (DFS). The puzzles explored will be of different difficulties for regular 9x9 grids and bigger 16x16 and 25x25 grids. This way the paper analyzes whether GAs are a viable option for efficiently solving Sudoku boards compared to naive search algorithms as the complexity of the puzzle increases.

This document is organized in the following way. Chapter one describes the objective of this study. Chapter two describes the problem of solving Sudoku puzzles and the motivation of using evolutionary algorithm approaches to solve this type of problem. Chapter three details the algorithm used in this study. The following chapter lays out the setup of the experiments run in this study. Lastly, chapter five presents and analyses the results of this paper.

2 Sudoku Puzzle Problem

2.1 Problem Description

Sudoku is a Japanese logical game that is played on a 9×9 grid, that are further divided into 3×3 subgrids. The objective of the game is to fill the grid with digits from 1 to 9, ensuring that each row, column, and subgrid contains each digit exactly once[3]. The puzzle starts with some cells already filled in, and the player must use logic and respect the base rule of the game in order to finish the puzzle.

Sudoku puzzles can vary in difficulty based on how many numbers they start with, the arrangement of these numbers and even the varying size of the sudoku, since they can go beyond the standard 9×9 grid, similar to the 25×25 grid we used to test our algorithm with. Thus we can think of the sudoku problem as a graph coloring problem if it were to be expressed in a mathematical context. The 9×9 grid can be seen as graph that has 81 vertices, namely one vertex for each cell. Each vertex can be labeled with an ordered pair (x,y) , where x and y are integers between 1 and 9 [4]. Two distinct vertices $(x1,y1)$ and $(x2,y2)$ are connected by an edge if and only if they are in the same row, column or subgrid:

- $x1 = x2$, which translates as same row
- $y1 = y2$, which translates as same column
- $\lfloor (x1 - 1)/3 \rfloor = \lfloor (x2 - 1)/3 \rfloor$ and $\lfloor (y1 - 1)/3 \rfloor = \lfloor (y2 - 1)/3 \rfloor$ which translates as same subgrid

As previously mentioned, the puzzle is completed when all the vertices have an integer between 1 and 9 assigned to them, in such a way that vertices that are joined by an edge don't have the same number assigned to them.

2.2 Evolutionary Approach

While Sudoku has a deterministic solution, evolutionary algorithms can be used in order to compare their performance with more traditional approaches such as the depth-first search algorithm. As it is mentioned in[5], they managed to implement an evolutionary algorithm that could solve very efficiently easy Sudoku puzzles. The downside of their algorithm, was that it wasn't as efficient when it came down to solving medium or hard puzzles. The way, they approached the algorithm [5] was to test constraints on different sets of spaces, namely the Hamming space and the Row Swap space. The drawn conclusions were that even though the Hamming

space had a smoother fitness landscape, in the end the Row Swap space was more efficient and produced more optimal results.

Another approach was taken in the paper [6], where pre-processing was involved. The pre-processing involved filling some sure numbers in the sudoku, through different methods, such as the naked single method, hidden single method, full house and lone rangers. The way in which a method was selected to be used was based on the type of puzzle, since each method has different efficiency depending on the case. After the pre-processing was done, different evolutionary algorithms were tested, such as a genetic algorithm that uses crossover and mutation operators, another one that firstly pre-processes the board and then applies the said genetic algorithm and finally even a genetic algorithm that uses an ant colony optimization. The fitness function used in this paper is made of three components: first one minimizes missing digits into rows and columns, the second one uses a sort of aging penalty, in order to prevent the same best individual persisting over generations. Finally, the third one penalizes violations where a candidate digit conflicts with a given clue in its row or column. The conclusion they arrived to was that solving the sudoku puzzle was faster when pre-processing was involved and that a hybrid algorithm between genetic algorithm and ant colony optimization was the most efficient with every type of sudoku puzzle.

Lastly, the paper that influenced our approach the most was [3], which uses the genetic algorithm with Darwinian evolution as its inspiration. This implies that each individual in the population is tested against a fitness function, that further decides whether the individual should be removed from the population or used as a parent to get closer to the desired result. New individuals are created through the use of crossover and mutation. Mutations is restricted to sub-blocks, using swap mutation as the primary operator. Further on, it was observed in the paper that using a cataclysmic mutation, specifically a restart of the population, every 2000 generation brought the best results when no solution was found. Similar to our project, the authors experimented with different fitness functions, but ended up using a function that counts missing or duplicate digits in each row and column, assigning penalties. This being said, the optimal value of the fitness function would be 0. On top of that, the authors introduced an aging penalty, adding +1 to the fitness of the best solution if it remains unchanged. The conclusion the paper arrived to was that the genetic algorithm could solve sudoku puzzles quite efficiently, and even if there were better performing algorithms, they sometimes would fail at puzzles that the genetic algorithm could solve.

The motivation for our project is quite straightforward. We want to explore if a streamlined evolutionary algorithm with carefully chosen operators can achieve

better performances than already existing and efficient search algorithms, such as the depth-first search algorithm. As previously mentioned our main inspiration for the evolutionary algorithm was the paper [3], thus we implemented an evolutionary algorithm that uses mutation and crossover operators, while also taking an elitist approach regarding the selection of the parents. On top of that, we also wanted to investigate the scalability of our evolutionary algorithm, thus we also tested it on a 16×16 puzzle. Evolutionary algorithms are well suited for this type of problem, because they can handle large solution spaces effectively, balance exploration and exploitation, and adapt by integrating problem-specific heuristics. All being considered, we believe that the Sudoku problem is a great choice for analyzing all those key aspects of evolutionary algorithms.

3 Genetic Algorithm

This chapter describes the implementations of two algorithms used to solve Sudoku puzzles. It starts out by briefly introducing depth-first search as the baseline measure, and then goes on to explain how each of the genetic algorithms work.

3.1 Baseline

Depth-first search (DFS) is used as a deterministic baseline. It iteratively scans the board for the first empty cell, attempts candidate values $1 \dots n$, and backtracks to try alternative values if no candidate leads to a solution. This approach ensures a solution (given that the puzzle is solvable) and is simple to implement, which makes it a reliable reference for correctness and for measuring runtimes.

Algorithm 1 shows the implementation of the DFS algorithm in pseudo-code.

Algorithm 1 Depth-first search algorithm

Require: matrix

Ensure: **true** if solved, **false** otherwise

```

 $n \leftarrow \text{length}(\text{matrix})$ 
for  $i \leftarrow 1$  to  $n$  do
  for  $j \leftarrow 1$  to  $n$  do
    if  $\text{matrix}[i, j] = 0$  then
      for  $\text{num} \leftarrow 1$  to  $n$  do
        if  $\text{is\_valid\_move}(\text{matrix}, i, j, \text{num})$  then
           $\text{matrix}[i, j] \leftarrow \text{num}$ 
          if  $\text{DFS}(\text{matrix})$  then
            return true
          end if
           $\text{matrix}[i, j] \leftarrow 0$ 
        end if
      end for
    end if
  end for
  return false
end if
end for
return true

```

3.2 Implementations

This section presents and compares two genetic algorithms designed to solve Sudoku puzzles. Each implementation varies in its approach to selection, crossover, mutation strategies, and fitness evaluation.

The first GA implementation originated from initial project discussions and represents our first practical attempt to apply standard genetic operators to Sudoku. The second implementation builds on that design and performs better in general, which is why we use it as the representative GA in the following chapters.

Nevertheless, we decided not to omit the first GA from this paper, because it represents the iterative approach to finding an optimized solution, and it shows which algorithms tend to not work as good for Sudoku compared to the algorithms in the latter implementation.

The following sections document these algorithms and use the Sudoku board shown in 1 as a reference for visualizing the selection and mutation step.

	2					3	1
7					3		
			1	4		2	9
	5	2	7	6	4		1
	6	3		1	2	7	5
	7	8				4	
2			3	7			5
	1					9	
5	4			8	1		

Figure 1: Initial Sudoku puzzle

3.2.1 Implementation 1

Algorithm 2 initializes a population by randomly filling empty cells and then iteratively evolves it using selection, per-cell crossover and mutation for up to a fixed number of generations. Most offspring are produced by top parents of the previous generation. Then, the population is replaced by the offspring, and eventually the best individual is returned.

A detailed description of how the fitness function, crossover, and mutation work, will be given afterwards.

Algorithm 2 Genetic Algorithm 1

Require: *initial_sudoku*, *population_size*, *mutation_rate*, *max_generations***Ensure:** Best individual found

```

population  $\leftarrow$  PopulateRandomly(initial_sudoku, population_size)
generation  $\leftarrow$  0
while generation < max_generations do
  population  $\leftarrow$  SortByFitness(population) ▷ descending order
  if Fitness(first(population)) = max_fitness then
    return first(population)
  end if
  next_generation  $\leftarrow$   $\emptyset$ 
  for i  $\leftarrow$  1 to population_size  $\times$  0.8 do
    parent1, parent2  $\leftarrow$  Sample(top 20% of population)
    child  $\leftarrow$  Crossover(parent1, parent2)
    child  $\leftarrow$  Mutate(child, mutation_rate, initial_sudoku)
    next_generation  $\leftarrow$  next_generation  $\cup$  {child}
  end for
  for i  $\leftarrow$  1 to population_size  $\times$  0.2 do
    parent1, parent2  $\leftarrow$  Sample(population)
    child  $\leftarrow$  Crossover(parent1, parent2)
    child  $\leftarrow$  Mutate(child, mutation_rate, initial_sudoku)
    next_generation  $\leftarrow$  next_generation  $\cup$  {child}
  end for
  population  $\leftarrow$  next_generation
  generation  $\leftarrow$  generation + 1
end while
population  $\leftarrow$  SortByFitness(population) ▷ descending order
return first(population)

```

Fitness function Let R_i , C_j and B_b denote the sets of distinct, non-zero digits present in row i , column j , and block b , respectively. Using zero-based indices ($i, j \in \{0, \dots, 8\}$) the block index is

$$b = 3 \left\lfloor \frac{i}{3} \right\rfloor + \left\lfloor \frac{j}{3} \right\rfloor.$$

The fitness of a board is the total number of distinct digits across all rows, columns and 3×3 blocks:

$$F = \sum_{i=0}^8 |R_i| + \sum_{j=0}^8 |C_j| + \sum_{b=0}^8 |B_b|.$$

This definition results in a maximum fitness score of 243 for 9×9 boards.

Crossover A new individual is creating by crossing over two parents. The crossover process simply iterates over all the tiles that were not given in the initial puzzle and takes the value for that tile from either parent to copy it to the child.

6	2	9	5	7	6	7	3	1		8	2	6	7	7	5	7	3	1		8	2	9	7	7	6	7	3	1
7	8	5	6	9	3	1	2	7		7	1	3	9	8	3	8	2	4		7	1	5	6	8	3	1	2	4
1	3	4	1	4	7	2	9	5		1	5	5	1	4	7	2	9	3		1	3	5	1	4	7	2	9	5
4	5	2	7	6	4	4	1	8		9	5	2	7	6	4	1	1	8		9	5	2	7	6	4	1	1	8
9	6	3	8	1	2	7	5	9		4	6	3	9	1	2	7	5	9		9	6	3	8	1	2	7	5	9
9	7	8	2	9	5	4	6	2		3	7	8	5	9	1	4	6	3		3	7	8	5	9	1	4	6	3
2	2	6	3	7	7	9	4	5		2	8	7	3	7	1	7	4	5		2	2	6	3	7	7	7	4	5
4	1	6	2	2	3	9	7	9		3	1	6	4	2	3	9	7	2		4	1	6	4	2	3	9	7	9
5	4	1	2	8	1	3	4	8		5	4	4	2	8	1	3	4	3		5	4	1	2	8	1	3	4	3

Figure 2: Exemplary selection step according to Implementation 1

In the figure above, two parents (left and middle) of the current generation form an individual of the next generation (right). One area for improvement is evident in row 7: Parent 1 contains two occurrences of the digit 7 and Parent 2 contains three. The offspring also has three 7s, while it would have been possible to generate a row with all distinct digits, namely **286371945**.

Mutation After the creation of a new individual, it undergoes up to *mutation_amount* mutations. A mutation means that a random tile changes its value. If that tile, however, happens to be a given tile from the initial puzzle, it is not altered.

8	2	9	7	7	6	7	3	1		8	2	9	7	7	6	7	3	1	
7	1	5	6	8	3	1	2	4		7	1	5	6	8	3	1	2	4	
1	3	5	1	4	7	2	9	5		1	3	5	1	4	7	2	9	5	
9	5	2	7	6	4	1	1	8		9	5	2	7	6	4	1	1	8	
9	6	3	8	1	2	7	5	9		9	6	3	8	1	2	7	5	9	
3	7	8	5	9	1	4	6	3		3	7	8	5	9	1	4	6	3	
2	2	6	3	7	7	7	4	5		2	2	6	3	7	7	7	4	5	
4	1	6	4	2	3	9	7	9		4	1	6	4	2	3	9	7	9	
5	4	1	2	8	1	3	4	3		5	4	1	2	8	1	3	4	3	

Figure 3: Exemplary mutation step according to Implementation 1

Figure 3 illustrates a mutation step with *mutation_amount* = 2. Meanwhile, only a single tile (row 7, column 7) was actually modified to have the value 7. The other candidate position coincided with a given cell and therefore could not be changed. This example highlights a weakness of the implementation: mutations can be wasted when selected positions are immutable, and a single mutation may decrease fitness significantly by introducing additional conflicts (the number of 7's

increased to 4 in the respective row, to 2 in the respective column, and to 3 in the respective grid).

Problem This implementation leads to a situation where the diversity of the population decreases rapidly, resulting in convergence to suboptimal solutions. On many runs, the algorithm creates more than 100,000 generations before finding a valid solution. To counter this problem, we have tried to increase the mutation rate and the population size, as well as introducing new random individuals in each generation. However, these adjustments only led to marginal improvements in performance.

3.2.2 Implementation 2

Algorithm 3 creates a population with unique rows and improves it using elitism, tournament selection, row-wise crossover and swap-based mutation. Each generation the population is evaluated, the best *elite_size* individuals are carried over unchanged, and the remaining offspring are produced from parents chosen by tournament; children copy whole rows from parents to preserve row validity and then undergo mutations that swap non-given tiles within rows. The algorithm terminates when a perfect solution is found or the maximum number of generations is reached.

A description of the fitness function, crossover, and mutation will once again be given afterwards.

Algorithm 3 Genetic Algorithm 2

Require: *population_size*, *elite_size*, *max_generations***Ensure:** Best individual found

```

population  $\leftarrow$  PopulateWithUniqueRows(population_size)
best_individual  $\leftarrow$  arbitrary element of population
for generation  $\leftarrow$  1 to max_generations do
    for all ind  $\in$  population do
        ComputeFitness(ind)
    end for
    population  $\leftarrow$  SortByFitness(population)  $\triangleright$  ascending order
    if Fitness(first(population)) < Fitness(best_individual) then
        best_individual  $\leftarrow$  first(population)
    end if
    if Fitness(best_individual) = 0 then
        return best_individual
    end if
    next_generation  $\leftarrow$  Top(elite_size, population)
    parents  $\leftarrow$  TournamentSelection(population)
    for i  $\leftarrow$  1 to population_size - elite_size step 2 do
        parent1  $\leftarrow$  parents[i]
        parent2  $\leftarrow$  parents[i + 1]
        child1  $\leftarrow$  Crossover(parent1, parent2)
        next_generation  $\leftarrow$  next_generation  $\cup$  {child1}
        if |next_generation| < population_size then
            child2  $\leftarrow$  Crossover(parent2, parent1)
            next_generation  $\leftarrow$  next_generation  $\cup$  {child2}
        end if
    end for
    population  $\leftarrow$  next_generation
end for
return best_individual

```

Fitness function Let C_j and B_b denote the sets of distinct, non-zero digits in column j and block b , respectively, and let n be the board size. The fitness function counts the number of duplicates in each column, and 3x3 subgrid:

$$F = \sum_{j=0}^{n-1} (n - |C_j|) + \sum_{b=0}^{n-1} (n - |B_b|).$$

A value $F = 0$ indicates a valid solution.

Crossover The crossover step is done by iterating through the rows of the sudoku board and copying each row from either parent. This ensures that the children will still respect the uniqueness of numbers in each row.

9	2	4	6	7	8	5	3	1
7	9	1	5	2	3	8	6	4
8	3	5	1	4	6	2	9	7
3	5	2	7	6	4	9	1	8
4	6	3	8	1	2	7	5	9
6	7	8	5	3	9	4	1	2
2	9	6	3	7	4	1	8	5
3	1	7	2	6	5	9	8	4
5	4	9	6	8	1	3	2	7

 $+$

8	2	4	7	9	6	5	3	1
7	9	1	5	2	3	6	8	4
3	6	7	1	4	8	2	9	5
3	5	2	7	6	4	9	1	8
8	6	3	4	1	2	7	5	9
1	7	8	5	3	9	4	2	6
2	9	6	3	7	4	1	8	5
3	1	7	8	5	6	9	4	2
5	4	6	9	8	1	3	2	7

 $=$

8	2	4	7	9	6	5	3	1
7	9	1	5	2	3	6	8	4
8	3	5	1	4	6	2	9	7
3	5	2	7	6	4	9	1	8
8	6	3	4	1	2	7	5	9
1	7	8	5	3	9	4	2	6
2	9	6	3	7	4	1	8	5
3	1	7	2	6	5	9	8	4
5	4	6	9	8	1	3	2	7

Figure 4: Exemplary selection step according to Implementation 2

In figure 4, it is easy to see that the very first row is taken from the first parent (left), while row 8 is copied from the second parent (center). Duplicate numbers, as a result of the crossover step, can now only emerge in columns or subgrids, which is an improvement compared to [Implementation 1](#).

Mutation For the mutation, we iterate all rows of an individual and swap two random tiles that both haven't been given in the initial sudoku. We do this with the probability of *mutation_rate* for each row, once again ensuring the uniqueness of numbers in each row.

8	2	4	7	9	6	5	3	1
7	9	1	5	2	3	6	8	4
8	3	5	1	4	6	2	9	7
3	5	2	7	6	4	9	1	8
8	6	3	4	1	2	7	5	9
1	7	8	5	3	9	4	2	6
2	9	6	3	7	4	1	8	5
3	1	7	2	6	5	9	8	4
5	4	6	9	8	1	3	2	7

 $\xrightarrow{\text{mutation}}$

8	2	4	6	9	7	5	3	1
7	9	1	5	2	3	6	8	4
8	3	5	1	4	6	2	9	7
3	5	2	7	6	4	9	1	8
4	6	3	8	1	2	7	5	9
1	7	8	5	3	9	4	2	6
2	9	6	3	7	4	1	8	5
3	1	7	2	6	5	9	8	4
5	4	6	9	8	1	3	2	7

Figure 5: Exemplary mutation step according to Implementation 2

The mutation of the individual that resulted from the previous crossover is shown in figure 5. For this example, the mutation rate was set to 20%, leading to mutations in row 1 and row 5.

Problem The second genetic algorithm also often converges to local optima and loses diversity, which limits further improvement. We added adaptive mutation rates and replaced most of the population when fitness stagnated across multiple generations. These measures helped but did not completely eliminate local optima problems. Furthermore, scalability is an issue: larger boards increase the combinatorial search space and therefore require proportionally larger populations and/or more generations to retain a comparable probability of finding a solution.

Despite these limitations, Implementation 2 performs much better in practice than Implementation 1, but the problems with premature convergence apply to both implementations and to genetic algorithms in general.

3.3 Solution representation

The solution representation is an $n \times n$ matrix where each entry is a number in the range from 1 to n . The solution is valid if and only if every number in a row, a column, and a respective block is unique.

4 Experimental part

This section describes the setup of experiments [?]:

- Provide the details of the hardware and software that you used.
- Describe the steps you carried out during your experiments.
- Detail the data you used for the evaluation of your algorithm.

This section describes the setup of the experiments run to study the efficiency of solving Sudoku puzzles with GA approaches.

The test will be run on a Macbook pro M1 max.

4.1 Choosing hyperparameters

Chapter 3 discusses the algorithm used for the experiments in detail. However, before the algorithm can be run, the hyperparameters for *population_size* and *mutation_rate* need to be chosen. The solution quality of a stochastic algorithm strongly depends on choosing the correct hyperparameters. Previous studies show that configuring hyperparameters of GAs using Bayesian Optimization leads to significantly better results than choosing them at random while keeping the computational time low[7].

Using Bayesian Optimization, the optimal values for the hyperparameters *population_size* and *mutation_rate* of the GA are evaluated as follows. First a search space is defined. It specifies the lower and upper bounds of the hyperparameters. Next, an *objective function* tries to solve a Sudoku using the GA. It measures solving time and the final fitness. These measurements calculate a score. An *optimize* method tries to minimize the score by fitting a Gaussian Process model to predict performance across the parameter space. It iteratively chooses new parameter sets that are either promising or not well explored. This process is repeated for a set number of iterations. The result is the best found population size and mutation rate.

4.2 Running the experiment

The GA is run 1000 times for a maximum of 10.000 generations for Sudoku puzzles of easy, medium and hard difficulty of 9x9 boards. Because of the drastic jump in complexity with 16x16 and 25x25 boards, the GA is only run ? times for these puzzles. For the evaluation of the performance the following properties of the experiment are tracked:

1. Number of successfully solved boards
2. Execution time
3. Number of average generations
4. Average execution time per Sudoku puzzle
5. Average generations per second

For failed runs the following properties are tracked:

1. Number of generations stuck at local minimum
2. Best fitness achieved
3. Average violations
4. Generations without improvement
5. Average generations without improvement

To compare the performance of the GA, a DFS algorithm is run on Sudoku boards of the same complexity. For the evaluation of the performance the following properties of the experiment are tracked:

1. Number of successfully solved boards
2. Execution time
3. Average execution time per Sudoku puzzle

4.3 Sudoku boards used in the experiment

Because of the high number of Sudoku boards needed to accurately evaluate the performance of the stochastic algorithm, the Sudoku grids are generated randomly as part of the experiment. The boards are generated in the following way. First, a solved board is generated using the DFS algorithm. Next, a random number is removed from the grid. The DFS algorithm tries to solve the puzzle again. If it is still solvable, repeat the previous steps until the requirement of given numbers is met for the different complexities of the Sudoku boards. For 9x9 grids, the easy difficulty gives between 40-50 numbers. Medium and hard puzzles have 30-39 and 25-29 givens respectively.

Because of the increased complexity, run time and reduced success rate of solving 16x16 and 25x25 puzzles, the algorithm will be run fewer times only on three 16x16 boards taken from [8] and one 25x25 board taken from [9].

5 Results and Analysis

5.1 Bayesian optimization results

We Bayesian optimization to find the optimal population size and mutation rate for the GA. The results are shown in Figure 6.

5.1.1 Optimization progress

From the first chart, the difference between the best score and the worst score is small. It means that the GA is stable and the hyperparameters are good.

5.1.2 Population size vs Performance

From the second chart, we can see population size around 100 and 900 can find the solution in a short time, while score is more discreted at 900 population size. Thus the average time of solving puzzle at 900 population size is costlier than 150 population size.

5.1.3 Mutation rate vs Performance

From the third chart, we can see muatation rate between 0.05 and 0.15 have many times not find the solution, however the score of failer cases become lower and no solution events happened less frequently when mutation rate pass 0.20.

5.1.4 Parameter space exploration

From the fourth chart, the best population size is 142 and the best mutation rate is 0.230. This result is based on the best score shows in first chart.




Figure 6: Bayesian optimization results.

5.2 Test the performance of the GA with varied difficulty

We test easy, medium and hard difficulty in same setting, and it's costly to test 1000 times for each difficulty. Thus we only test easy difficulty for 1000 times, medium one for around 400 times, and hard one for 100 times.

5.2.1 Easy difficulty

The charts below shows the performance of the GA with easy difficulty. There are more than 500 times find solution in less than 0.5 seconds. only around 1/20 solutions were reached very slow which is more than 20s. Additionally, the average execution time is 5s-ish, which is a good performance. Moreover, the execution time and generation have linear relationship. It means that if algorithm need more generations to get the solution, the more execution time it will cost.



Figure 7: Generation vs execution time for easy difficulty puzzles.




Figure 8: Generation and execution time distribution for easy difficulty puzzles.

5.2.2 Medium difficulty

The charts below shows the performance of the GA with medium difficulty, the time cost of founding solution at quick and medium time only play a small role in total cases. However, the proportion of the getting solution in very slow is 60%,

which is much higher comparing with easy difficulty. Meanwhile, the execution time and generation still remain a linear relationship.




Figure 9: Generation vs execution time for medium difficulty puzzles.




Figure 10: Generation and execution time distribution for medium difficulty puzzles.

5.2.3 Hard difficulty

The charts below shows the performance of the GA with hard difficulty, the proportion of the getting solution in very slow is 97%, which is much higher comparing with easy and medium difficulty. Meanwhile, the execution time and generation still remain a linear relationship.




Figure 11: Generation vs execution time for hard difficulty puzzles.

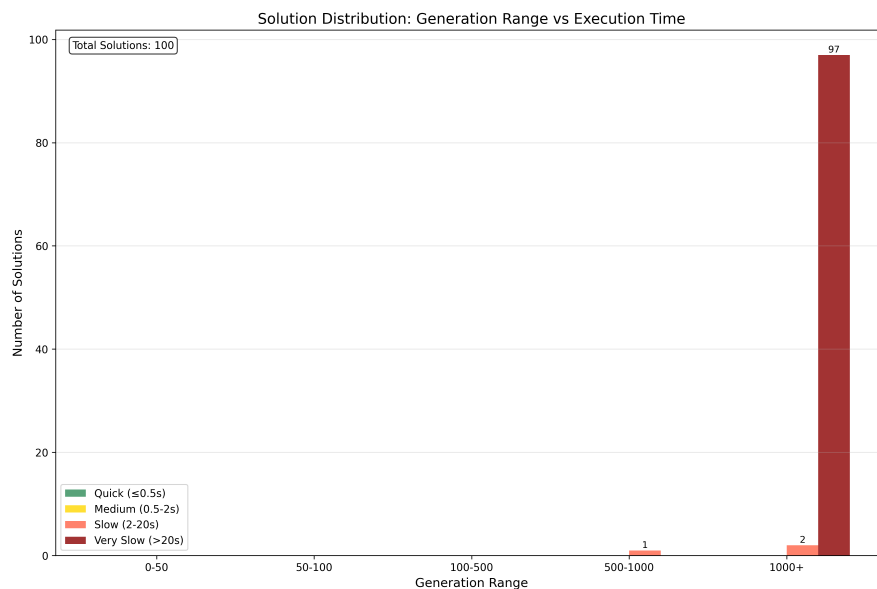


Figure 12: Generation and execution time distribution for hard difficulty puzzles.

5.3 Test the performance by rule based algorithm

We test the performance of the rule based algorithm with easy, medium and hard difficulty.




Figure 13: sudoku_dfs_log_easy_histogram.



Figure 14: sudoku_dfs_log_medium_histogram.




Figure 15: sudoku_dfs_log_hard_histogram.

The charts illustrate that no matter how hard the puzzle is, the rule based algorithm can solve it in a short time. Moreover, the execution time is very stable, only takes 0.001s-0.002s to solve each puzzle.

6 Conclusions

In this section you should provide a concise summary of what has been done, the obtained results and some recommendations on how this study could be extended.

This paper studied the efficiency of solving Sudoku puzzles with GA approaches of different complexities. The results show that the performance on 9x9 boards varies drastically. While most easy boards can be solved fast, the algorithm struggles on medium and hard boards. Only 14% of medium boards can be solved in under 2 second. For puzzles of hard difficulties the algorithm did not achieve a single solve in under 2 seconds. Additionally, the success rate of the GA in finding a solution to the puzzle decreased with increasing complexity of the problem. While 100% of the easy boards were solved, only x% and y% of boards were solved of medium and hard complexity respectively.

The naive search algorithm DFS found solutions for every puzzle of every difficulty. Furthermore, the mean solving time of the boards of hard difficulty is 0.112 seconds, while easier puzzles were solved even faster.

These results show that while GAs can be used to solve Sudoku puzzles, they do not outperform naive search approaches. This is the case in the Sudoku context because there is only one optimal solution (solving the board without any duplicates). GAs are not optimal. They perform well in finding solutions that are "good enough". When solving Sudokus however, only the optimal solution is relevant. The algorithm finds local maxima quickly, but then gets stuck in those in many cases. The algorithm reacts by soft resetting the population. This is done by generating a new population and combining it with some of the worst performing individuals. This helps the algorithm to get out of a local maximum, but I doesn't guarantee to find the optimal solution next. This is the reason why y% of hard difficulty boards could not be solved.

This led to the decision of not evaluating the performance of the bigger 16x16 and 25x25 boards in this paper. It is expected that the performance of bigger grids is even worse, especially because the increase in complexity is far greater when increasing the size of the boards compared reducing given numbers in a 9x9 puzzle.

Future research may consist of combining GAs with different optimization approaches. If a significant increase in performance on 9x9 boards can be achieved, these hybrid GAs may be suitable for solving bigger 16x16 and 25x25 Sudoku puzzles.

As shown in [?], evolutionary algorithms can solve Sudoku puzzles efficiently.

References

- [1] “Sudoku - wikipedia.” Accessed: 2025-10-01.
- [2] “Genetic algorithm - wikipedia.” Accessed: 2025-10-01.
- [3] T. Mantere and J. Koljonen, “Solving, rating and generating sudoku puzzles with ga,” *2007 IEEE Congress on Evolutionary Computation, CEC 2007*, pp. 1382–1389, 2007.
- [4] “Mathematics - wikipedia.”
- [5] “Evolutionary algorithms and sudoku.”
- [6] P. R. Amil and T. Mantere, “Solving sudoku’s by evolutionary algorithms with pre-processing,” *Advances in Intelligent Systems and Computing*, vol. 837, pp. 3–15, 2019.
- [7] C. R  ther and J. Rieck, “A bayesian optimization approach for tuning a genetic algorithm solving practical-oriented pickup and delivery problems,”
- [8] “Free 16x16 printable sudoku puzzles (pdf).”
- [9] “The daily printable 25 x 25 sudoku puzzle.”