

Group 6 Assignment Report

An Evolutionary Algorithm for solving Sudokus

Ben Ertugrul Eduard Iorga Haoran Liu
Moritz Schlitter

September 29, 2025

1 Introduction

The recommended software for typesetting assignment reports is L^AT_EX. It will allow you to prepare high-quality documents, especially in the area of Computer Science. This document can serve as a template for reports. Each section begins with brief instructions in red text. All the instructions in red, as well as the dummy text, should be removed in the final version to submit. The L^AT_EX source of this file includes examples of using the most needed commands and environments. You can find plenty of other examples with explanations in many web forums and discussion groups on the Internet. The easiest way to edit your report is to use <https://www.overleaf.com/>. Overleaf does not require any setup on your computer, and it is free to create an account.

The book *Writing for Computer Science* [?] is a useful assistance on how to write properly and present your work when it comes to Computer Science topics. It is a strong recommendation to follow its guidelines and limit the usage of AI tools to generate text. Keep in mind that the examiner is an expert in Evolutionary Computation and therefore, any false information generated by an AI tool is easily notable. Such case may lead to failing the assignment.

The introduction should briefly introduce the assignment and its purpose.

-
- Context - Introduce Sudoku and Genetic algorithm very shortly
 - This paper studies the efficiency of solving Sudoku puzzles with evolutionary algorithm approaches
 - Different difficulties for 9x9, increased sizes
 - Additionally, we want to test whether increasing the size of the Sudoku board leads to better performance of evolutionary algorithms compared to naive search algorithms like DFS.
 - The performance is compared between evolutionary algorithms of different complexities as well as to Depth-First Search - a naive optimization algorithm
 - The objective is to draw a conclusion on which approach solves Sudoku more efficiently, both for regular 9x9 Sudoku boards as well as for bigger boards (16x16 and 25x25).
 - The document is structured as follows: Section 1: Introduction and goal of the paper. Section 2: Describes the problem the paper solves in detail (Solving Sudoku) + motivation of the evolutionary approach used. Section 3: Describes the algorithm

in detail. Section 4: Describes the experiments run with the algorithm. Section 5: Presents and analyses the results. Section 6: Conclusion.

Traditionally, a Sudoku is a logic based number placement puzzle. The objective of the puzzle is to fill a most commonly 9x9 sized grid with digits so that each column, row and 3x3 subgrid contain all of the digits from 1 to 9.[1]

Solving such a puzzle programmatically falls into the category of search and optimization problems. These types of problems can be approached in different ways. One possible approach is depicted by evolutionary algorithms or in this case more specifically, genetic algorithms (GA).[1]

Inspired by the process of natural selection, GAs use biologically inspired operators such as selection, crossover and mutation to generate solutions to optimization and search problems.

This paper studies the efficiency of solving Sudoku puzzles with GA approaches. The objective is to compare the performance of different implementations of GAs. These approaches will additionally be compared to the naive search algorithm depth-first search (DFS). The puzzles explored will be of different difficulties for regular 9x9 grids and bigger 16x16 and 25x25 grids. This way the paper analyzes whether GAs are a viable option for efficiently solving Sudoku boards compared to naive search algorithms as the complexity of the puzzle increases.

This document is organized in the following way. Chapter one describes the objective of this study. Chapter two describes the problem of solving Sudoku puzzles and the motivation of using evolutionary algorithm approaches to solve this type of problem. Chapter three details the algorithm used in this study. The following chapter lays out the setup of the experiments run in this study. Lastly, chapter five present and analyses the results of this paper.

2 Problem this is the problem

The second section should present the problem you tackle using your evolutionary approach. Overall, this section should include:

- The mathematical formulation considered in your study. Some problems have a clear mathematical model (e.g., Travelling Salesman Problem), while others do not (e.g., n -Queens). Based on the problem you chose, search the literature and find a proper way to present the problem.
- One paragraph that briefly presents at least 3 published academic works where any evolutionary approach is used to solve the problem. It would be wise to cite here works that influenced your algorithm. This practice saves you time from looking for additional academic resources. You can find more information about reading and searching in the literature in [?].
- The motivation behind the evolutionary approach you decided to develop. A good practice would be to align the motivation with some literature gap found in the academic works you presented above. However, this is not mandatory. You can motivate your selection on the characteristics of the algorithm making it proper for the problem.

Note: Change the section's title to match the name of the problem you chose for your assignment.

-
- This section presents the problem this paper solves with the use of evolutionary algorithm approach.
 - Math in solving Sudokus? Maybe rather introducing quickly a sudoku, how to solve it and describing the steep increase of difficulty with increasing sizes of sudoku puzzles
 - Cite 3 published papers. Cite where our idea is from[2]. Another Paper[3].
 - Motivation behind evolutionary approach.

3 Algorithm

The third section should present the evolutionary approach you developed. You can divide this section into subsection. In any case, you should mention the following details:

Evolutionary approach. Clearly describe the algorithm you developed. You should clearly explain the evolutionary operators you used and what modifications you did to match the problem. It is extremely important to present also a pseudocode of your algorithm. An example is given in 1, below. For more insight into presentation of algorithms, you can advise [?].

To typeset pseudocode in L^AT_EX you can use one of the following options:

- Choose ONE of the (algpseudocode OR algcompatible OR algorithmic) packages to typeset algorithm bodies, and the algorithm package for captioning the algorithm.
- The algorithm2e package.

You can find more information here: <https://www.overleaf.com/learn/latex/Algorithms>

Algorithm 1 Example of an algorithm's pseudocode

Require: $n \geq 0$

Ensure: $y = x^n$

$y \leftarrow 1$

$X \leftarrow x$

$N \leftarrow n$

while $N \neq 0$ **do**

if N is even **then**

$X \leftarrow X \times X$

$N \leftarrow \frac{N}{2}$

▷ This is a comment

else if N is odd **then**

$y \leftarrow y \times X$

$N \leftarrow N - 1$

end if

end while

Solution representation. Clearly describe the solution representation you used. You can use figures to improve the comprehensibility of this part.

Fitness function. It is also very important to mention the fitness function you used. In many cases, the objective function of the problem is not the same as

the fitness function used in an evolutionary algorithm. An example, following the principles of [?], is given below.

$$F = \sum_{i=1}^d x_i^2 \quad (1)$$

where x_i is the i -th gene (i.e., decision variable) in the solution and d corresponds to the number of decision variables in the problem.

Note: Change the section's title to match the name of the algorithm you developed for your assignment.

3.1 Evolutionary Algorithm Implementations

This section presents and compares different implementations of an evolutionary algorithm designed to solve Sudoku puzzles. Each implementation varies in its approach to selection, crossover, mutation strategies, and fitness evaluation.

The following sections document these processes and use the Sudoku board shown in [Initial Sudoku puzzle](#) as a reference for visualizing the selection and mutation step.

	2					3	1
7					3		
			1	4		2	9
	5	2	7	6	4		1
	6	3		1	2	7	5
	7	8				4	
2			3	7			5
	1					9	
5	4			8	1		

Figure 1: Initial Sudoku puzzle

3.1.1 Implementation 1

1. Initialize the population with ***population_size*** random individuals that all respect the initially given sudoku tiles. These individuals will start out with random values in the empty tiles. The fitness of each individual is calculated as the sum of the number of unique numbers in each row, column, and 3x3

- subgrid. The maximum fitness is 243 (9 unique numbers * 9 rows + 9 columns + 9 subgrids).
- The next generation of individuals is created by selecting two parents (or the same one twice) randomly from the top 20% of the current population. These top parents will create 80% of the next generation's individuals. The other 20% of individuals will have parents that are chosen the entire population.
 - We perform random crossover, iterating over each cell and randomly choosing whether the child inherits the value from the first or second parent.
 - Afterwards, the child mutates *mutation_amount* random tiles on the board, changing their value. Any change on the initially given sudoku tiles is not allowed.
 - Once the new generation is created, we calculate the fitness of each individual and repeat the process from step 2. until a valid solution is found.

Selection The Exemplary selection step according to Implementation 1 depicts how two parents (left and middle) of the current generation form an individual of the next generation (right). One area for improvement is evident in row 7: Parent 1 contains two occurrences of the digit 7 and Parent 2 contains three. The offspring also has three 7s, while it would have been possible to generate a row with all distinct digits, namely **286371945**.

6	2	9	5	7	6	7	3	1		8	2	6	7	7	5	7	3	1		8	2	9	7	7	6	7	3	1
7	8	5	6	9	3	1	2	7		7	1	3	9	8	3	8	2	4		7	1	5	6	8	3	1	2	4
1	3	4	1	4	7	2	9	5		1	5	5	1	4	7	2	9	3		1	3	5	1	4	7	2	9	5
4	5	2	7	6	4	4	1	8		9	5	2	7	6	4	1	1	8		9	5	2	7	6	4	1	1	8
9	6	3	8	1	2	7	5	9		4	6	3	9	1	2	7	5	9		9	6	3	8	1	2	7	5	9
9	7	8	2	9	5	4	6	2		3	7	8	5	9	1	4	6	3		3	7	8	5	9	1	4	6	3
2	2	6	3	7	7	9	4	5		2	8	7	3	7	1	7	4	5		2	2	6	3	7	7	7	4	5
4	1	6	2	2	3	9	7	9		3	1	6	4	2	3	9	7	2		4	1	6	4	2	3	9	7	9
5	4	1	2	8	1	3	4	8		5	4	4	2	8	1	3	4	3		5	4	1	2	8	1	3	4	3

Figure 2: Exemplary selection step according to Implementation 1

Mutation Figure 3 illustrates a mutation step with mutation rate $n = 2$; however, only a single tile (row 7, column 7) was actually modified. The other candidate position coincided with a given cell and therefore could not be changed. This behaviour highlights a weakness of the implementation: mutations can be wasted when selected positions are immutable, and a single mutation may decrease fitness by introducing additional conflicts (in this example the row contains three 7s, the

column contains one 7 and the corresponding 3×3 block contains two 7s), which lowers the individual's score.

8	2	9	7	7	6	7	3	1
7	1	5	6	8	3	1	2	4
1	3	5	1	4	7	2	9	5
9	5	2	7	6	4	1	1	8
9	6	3	8	1	2	7	5	9
3	7	8	5	9	1	4	6	3
2	2	6	3	7	7	7	4	5
4	1	6	4	2	3	9	7	9
5	4	1	2	8	1	3	4	3

mutation
→

8	2	9	7	7	6	7	3	1
7	1	5	6	8	3	1	2	4
1	3	5	1	4	7	2	9	5
9	5	2	7	6	4	1	1	8
9	6	3	8	1	2	7	5	9
3	7	8	5	9	1	4	6	3
2	2	6	3	7	7	7	7	5
4	1	6	4	2	3	9	7	9
5	4	1	2	8	1	3	4	3

Figure 3: Exemplary mutation step according to Implementation 1

Problem This implementation leads to a situation where the diversity of the population decreases rapidly, resulting in premature convergence to suboptimal solutions. On most of the runs, the algorithm creates more than 100,000 generations before finding a valid solution. To counter this problem, we have tried to increase the mutation rate and the population size, as well as introducing new random individuals in each generation. However, these adjustments only led to marginal improvements in performance.

3.1.2 Implementation 2

1. Initialize the population with *population_size* number of random individuals that all respect the initially given sudoku tiles. The individuals create a complete board by iterating over the rows and filling in the empty tiles with random numbers that do not already exist in the respective row. The fitness function works counterintuitively, calculating the fitness as the sum of the number of duplicates in each column, and 3×3 subgrid. The minimum fitness is 0 (no duplicates).
2. The next generation starts with the top *elites_size* number of individuals from the current population. Then, a tournament selection is performed to select parents for the remaining individuals in the next generation. In each tournament, *K* individuals are randomly chosen from the population, and the one with the best fitness is selected as a parent. This selection process is repeated until we have *population_size* number of parents.
3. The crossover step is done by iterating over all parents in pairs, creating two children from each pair. For each row in the sudoku board, the child inherits

that row either from the first or the second parent. This ensures that the children will still respect the uniqueness of numbers in each row.

4. For mutation, we iterate all rows of an individual and swap two random tiles that both haven't been given in the initial sudoku. We do this with the probability of *mutation_rate* for each row, once again ensuring the uniqueness of numbers in each row.
5. Once the new generation is created, we calculate the fitness of each individual and repeat the process from step 2. until a valid solution is found.

```

+=====+=====+=====+ + +=====+=====+=====+ = +=====+=====+=====+ ~ +=====
| 9 2 4 | 6 7 8 | 5 3 1 | + | 8 2 4 | 7 9 6 | 5 3 1 | = | 8 2 4 | 7 9 6 | 5 3 1 | ~ | 8 2 4 | 6 9 7 |
| 7 9 1 | 5 2 3 | 8 6 4 | + | 7 9 1 | 5 2 3 | 6 8 4 | = | 7 9 1 | 5 2 3 | 6 8 4 | ~ | 7 9 1 | 5 2 3 |
| 8 3 5 | 1 4 6 | 2 9 7 | + | 3 6 7 | 1 4 8 | 2 9 5 | = | 8 3 5 | 1 4 6 | 2 9 7 | ~ | 8 3 5 | 1 4 6 |
+=====+=====+=====+ + +=====+=====+=====+ = +=====+=====+=====+ ~ +=====
| 3 5 2 | 7 6 4 | 9 1 8 | + | 3 5 2 | 7 6 4 | 9 1 8 | = | 3 5 2 | 7 6 4 | 9 1 8 | ~ | 3 5 2 | 7 6 4 |
| 4 6 3 | 8 1 2 | 7 5 9 | + | 8 6 3 | 4 1 2 | 7 5 9 | = | 8 6 3 | 4 1 2 | 7 5 9 | ~ | 4 6 3 | 8 1 2 |
| 6 7 8 | 5 3 9 | 4 1 2 | + | 1 7 8 | 5 3 9 | 4 2 6 | = | 1 7 8 | 5 3 9 | 4 2 6 | ~ | 1 7 8 | 5 3 9 |
+=====+=====+=====+ + +=====+=====+=====+ = +=====+=====+=====+ ~ +=====
| 2 9 6 | 3 7 4 | 1 8 5 | + | 2 9 6 | 3 7 4 | 1 8 5 | = | 2 9 6 | 3 7 4 | 1 8 5 | ~ | 2 9 6 | 3 7 4 |
| 3 1 7 | 2 6 5 | 9 8 4 | + | 3 1 7 | 8 5 6 | 9 4 2 | = | 3 1 7 | 2 6 5 | 9 8 4 | ~ | 3 1 7 | 2 6 5 |
| 5 4 9 | 6 8 1 | 3 2 7 | + | 5 4 6 | 9 8 1 | 3 2 7 | = | 5 4 6 | 9 8 1 | 3 2 7 | ~ | 5 4 6 | 9 8 1 |
+=====+=====+=====+ + +=====+=====+=====+ = +=====+=====+=====+ ~ +=====

```

4 Experimental part

This section describes the setup of experiments [?]:

- Provide the details of the hardware and software that you used.
- Describe the steps you carried out during your experiments.
- Detail the data you used for the evaluation of your algorithm.

Quisque ullamcorper placerat ipsum. Cras nibh. Morbi vel justo vitae lacus tincidunt ultrices. Lorem ipsum dolor sit amet, consectetur adipiscing elit. In hac habitasse platea dictumst. Integer tempus convallis augue. Etiam facilisis. Nunc elementum fermentum wisi. Aenean placerat. Ut imperdiet, enim sed gravida sollicitudin, felis odio placerat quam, ac pulvinar elit purus eget enim. Nunc vitae tortor. Proin tempus nibh sit amet nisl. Vivamus quis tortor vitae risus porta vehicula.

5 Results and Analysis

This section should present the obtained results and provide an insightful analysis of them. You can present the results using graphs, tables, or any other visualization method suits your purpose. Do not forget to include proper captions [?] in any of these illustration methods you use. You do not need to provide any execution details as they are already presented in Sec. 4.

A good practise would be to compare your algorithm with a simpler approach, such as (a) a naive method, (b) a Hill Climbing approach, or (c) a simple evolutionary algorithm. In the third case, you can use the simpler version of the algorithm you developed, i.e., the original algorithm without your modifications. In that case, you should briefly describe the comparing method(s) in Sec. 4. Alternatively, you can use some reference results derived from the repositories you found some benchmark instances.

To display tables, the `booktabs` package might be useful. For example, Table 1 shows how you should increase the size of n , when running your code. You can advice [?] to see a few examples of proper tables.

Table 1: Example of comparison the developed algorithm's results with the best ones from a repository.

Instance	Optimum (Repository xyz)	EA	time (s)
st70	678.597	677.109	0.67
ei176	545.387	544.369	1.16
kroA100	21285.443	21285.443	1.69
rd100	7910.396	7910.396	2.14
Pr136	96772	96770.924	7.11
Pr144	58537	58535.221	7.97
a280	2856.769	2856.769	33.47

You can use different illustration methods to present different aspects of your analysis. Figure 4 gives an example using the `pgfplots` package.

Quisque ullamcorper placerat ipsum. Cras nibh. Morbi vel justo vitae lacus tincidunt ultrices. Lorem ipsum dolor sit amet, consectetur adipiscing elit. In hac habitasse platea dictumst. Integer tempus convallis augue. Etiam facilisis. Nunc elementum fermentum wisi. Aenean placerat. Ut imperdiet, enim sed gravida sollicitudin, felis odio placerat quam, ac pulvinar elit purus eget enim. Nunc vitae tortor. Proin tempus nibh sit amet nisl. Vivamus quis tortor vitae risus porta vehicula.




Figure 4: Example of convergence analysis.

6 Conclusions

In this section you should provide a concise summary of what has been done, the obtained results and some recommendations on how this study could be extended.

Quisque ullamcorper placerat ipsum. Cras nibh. Morbi vel justo vitae lacus tincidunt ultrices. Lorem ipsum dolor sit amet, consectetur adipiscing elit. In hac habitasse platea dictumst. Integer tempus convallis augue. Etiam facilisis. Nunc elementum fermentum wisi. Aenean placerat. Ut imperdiet, enim sed gravida sollicitudin, felis odio placerat quam, ac pulvinar elit purus eget enim. Nunc vitae tortor. Proin tempus nibh sit amet nisl. Vivamus quis tortor vitae risus porta vehicula.

References

- [1] “Sudoku - wikipedia.”
- [2] T. Mantere and J. Koljonen, “Solving, rating and generating sudoku puzzles with ga,” *2007 IEEE Congress on Evolutionary Computation, CEC 2007*, pp. 1382–1389, 2007.
- [3] P. R. Amil and T. Mantere, “Solving sudoku’s by evolutionary algorithms with pre-processing,” *Advances in Intelligent Systems and Computing*, vol. 837, pp. 3–15, 2019.